

DIVING INTO THE MACHINE ROOM

Lecture 2

MAL2, Spring 2025



Today's Goals:

Understanding the how's and the why's of building a good neural network.

DIVING INTO THE MACHINE ROOM

- How training a neural network works
- Activation functions
- Faster optimizers
- Learning rate scheduling
- Regularization
- General suggestions

GRADIENT DESCENT

gradient

weights & biases

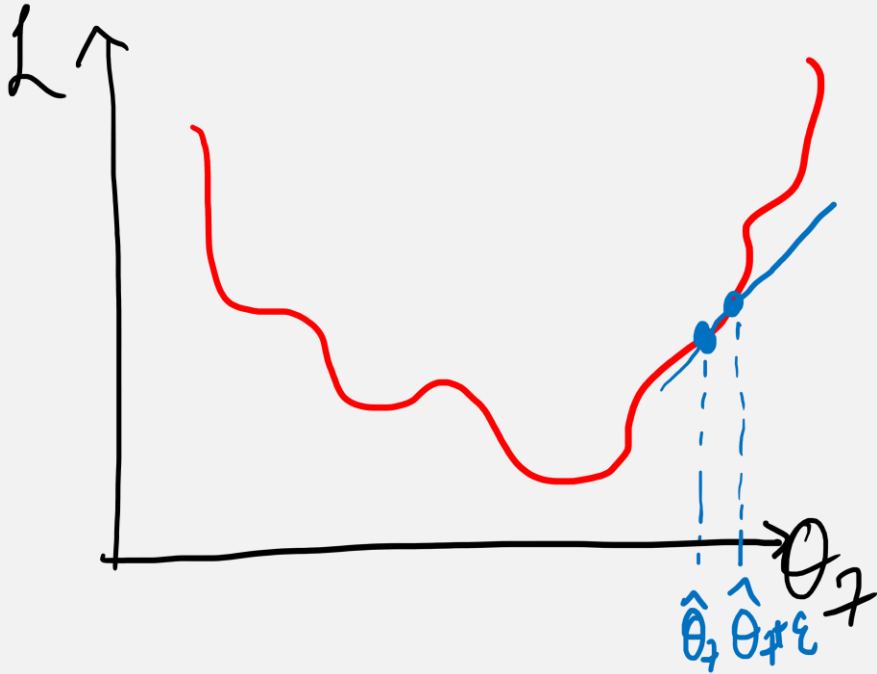
loss
function

$$\nabla L(\theta) = \left(\frac{\partial L}{\partial \theta_0}, \frac{\partial L}{\partial \theta_1}, \frac{\partial L}{\partial \theta_2}, \dots, \frac{\partial L}{\partial \theta_n} \right) \leftarrow \text{how do we calculate this?}$$

$$\theta_{\text{new}} = \theta_{\text{old}} - \eta \nabla L(\theta)$$

repeat until $\nabla L(\theta) \approx 0$

GRADIENT DESCENT – IDEA #1




$$\frac{\partial L}{\partial \theta_7} \approx \frac{L(\theta_7 + \epsilon) - L(\theta_7)}{\epsilon}$$

INTRACTABLE

This procedure must then be repeated for all parameters at every single training step ... and there may be **hundreds of thousands** of parameters!

GRADIENT DESCENT – IDEA #2

Use the chain rule



9 parameters
 L is optimized
 in 9D space

$$L = \text{MSE} \sim (\hat{y} - y)^2$$

we want $\frac{\partial L}{\partial w}$
 for all parameters
 in the model


$$\frac{\partial L}{\partial w_{21}^{(1)}}$$

It gets more complicated with larger networks and multiple "paths" to the same parameter – but the idea is there.

GRADIENT DESCENT – IDEA #2

$$\frac{\partial L}{\partial w_{21}^{(1)}}$$

$$L = (\hat{y} - y)^2$$



$$\frac{\partial L}{\partial w_{21}^{(1)}} = \frac{\partial L}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial a_1} \frac{\partial a_1}{\partial z_1} \frac{\partial z_1}{\partial w_{21}^{(1)}}$$

\downarrow \downarrow \downarrow \downarrow
 $2(\hat{y} - y)$ $w_{11}^{(2)}$ σ' x_2

$$\hat{y} = w_{11}^{(2)} a_1 + w_{21}^{(2)} a_2 + b_1^{(2)}$$

$$\sigma(z_1)$$

$$\sigma(z_2)$$


$$z_1 = w_{11}^{(1)} x_1 + w_{21}^{(1)} x_2 + b_1^{(1)}$$

$$z_2 = w_{12}^{(1)} x_1 + w_{22}^{(1)} x_2 + b_2^{(1)}$$

THE BACKPROPAGATION ALGORITHM

$$y=8$$

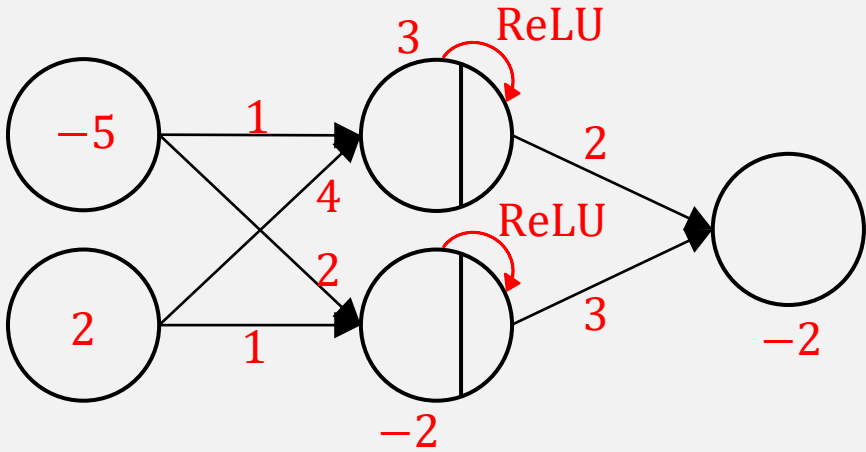
$$\hat{y}=10$$



$$\nabla L = \begin{pmatrix} \partial L / \partial w_{11}^{(1)} \\ \partial L / \partial w_{21}^{(1)} \\ \partial L / \partial b_1^{(1)} \\ \partial L / \partial w_{12}^{(1)} \\ \partial L / \partial w_{22}^{(1)} \\ \partial L / \partial b_2^{(1)} \\ \partial L / \partial w_{11}^{(2)} \\ \partial L / \partial w_{21}^{(2)} \\ \partial L / \partial b_1^{(2)} \end{pmatrix} = \begin{pmatrix} -40 \\ 16 \\ 8 \\ 0 \\ 0 \\ 0 \\ 24 \\ 0 \\ 4 \end{pmatrix}$$

$$\frac{\partial L}{\partial w_{21}^{(1)}} = 2(10 - 8) \cdot 2 \cdot 1 \cdot 2 = 16$$


THE BACKPROPAGATION ALGORITHM



$$\begin{pmatrix} w_{11}^{(1)} \\ w_{21}^{(1)} \\ b_1^{(1)} \\ w_{12}^{(1)} \\ w_{22}^{(1)} \\ b_2^{(1)} \\ w_{11}^{(2)} \\ w_{21}^{(2)} \\ b_1^{(2)} \end{pmatrix} = \begin{pmatrix} 1 \\ 4 \\ 3 \\ 2 \\ 1 \\ -2 \\ 2 \\ 3 \\ -2 \end{pmatrix} - 0.002 \begin{pmatrix} -40 \\ 16 \\ 8 \\ 0 \\ 0 \\ 0 \\ 24 \\ 0 \\ 4 \end{pmatrix} = \begin{pmatrix} 1.08 \\ 3.968 \\ 2.984 \\ 2 \\ 1 \\ -2 \\ 1.952 \\ 3 \\ -2.008 \end{pmatrix}$$

$$\nabla L = \begin{pmatrix} \partial L / \partial w_{11}^{(1)} \\ \partial L / \partial w_{21}^{(1)} \\ \partial L / \partial b_1^{(1)} \\ \partial L / \partial w_{12}^{(1)} \\ \partial L / \partial w_{22}^{(1)} \\ \partial L / \partial b_2^{(1)} \\ \partial L / \partial w_{11}^{(2)} \\ \partial L / \partial w_{21}^{(2)} \\ \partial L / \partial b_1^{(2)} \end{pmatrix} = \begin{pmatrix} -40 \\ 16 \\ 8 \\ 0 \\ 0 \\ 0 \\ 24 \\ 0 \\ 4 \end{pmatrix}$$

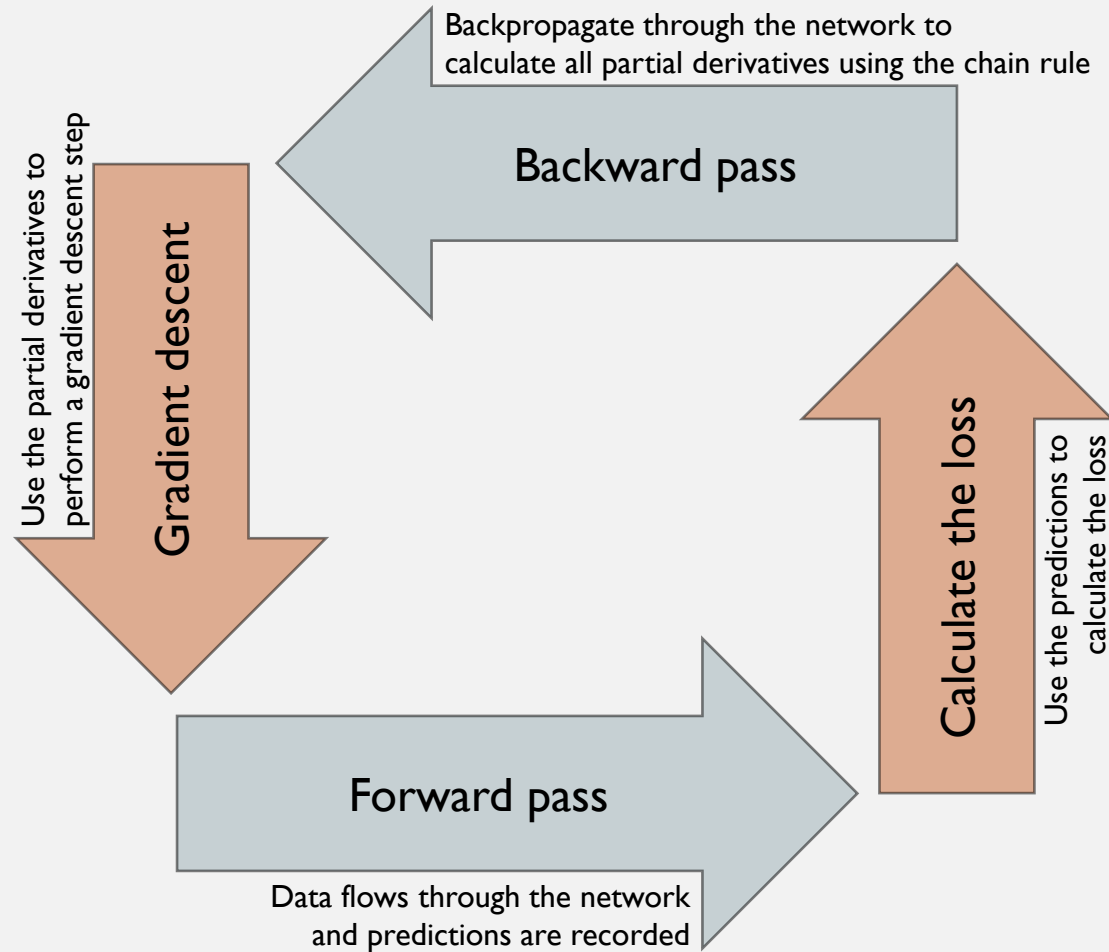
THE BACKPROPAGATION ALGORITHM



→ closer to 8
now repeat

$$\begin{pmatrix} w_{11}^{(1)} \\ w_{21}^{(1)} \\ b_1^{(1)} \\ w_{12}^{(1)} \\ w_{22}^{(1)} \\ b_2^{(1)} \\ w_{11}^{(2)} \\ w_{21}^{(2)} \\ b_1^{(2)} \end{pmatrix} = \begin{pmatrix} 1 \\ 4 \\ 3 \\ 2 \\ 1 \\ -2 \\ 2 \\ 3 \\ -2 \end{pmatrix} - 0.002 \begin{pmatrix} -40 \\ 16 \\ 8 \\ 0 \\ 0 \\ 0 \\ 24 \\ 0 \\ 4 \end{pmatrix} = \begin{pmatrix} 1.08 \\ 3.968 \\ 2.984 \\ 2 \\ 1 \\ -2 \\ 1.952 \\ 3 \\ -2.008 \end{pmatrix}$$

THE BACKPROPAGATION ALGORITHM




WHY DID WE JUST DO THIS?

$$\nabla L = \begin{pmatrix} \partial L / \partial w_{11}^{(1)} \\ \partial L / \partial w_{21}^{(1)} \\ \partial L / \partial b_1^{(1)} \\ \partial L / \partial w_{12}^{(1)} \\ \partial L / \partial w_{22}^{(1)} \\ \partial L / \partial b_2^{(1)} \\ \partial L / \partial w_{11}^{(2)} \\ \partial L / \partial w_{21}^{(2)} \\ \partial L / \partial b_1^{(2)} \end{pmatrix} = \begin{pmatrix} -40 \\ 16 \\ 8 \\ 0 \\ 0 \\ 0 \\ 24 \\ 0 \\ 4 \end{pmatrix}$$

-why are these numbers 0?

$$\frac{\partial L}{\partial w_{12}^{(1)}} = \frac{\partial L}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial a_2} \frac{\partial a_2}{\partial z_2} \frac{\partial z_2}{\partial w_{12}^{(1)}} = 2(\hat{y} - y) w_{21}^{(2)} \sigma' x_1 = 0$$

derivative of ReLU




"dying ReLU"

keeps outputting zero, killing the neuron

one of many problems we face →

In general: Gradients are unstable in deep neural networks!

VANISHING & EXPLODING GRADIENTS



gradients tend to
get smaller & smaller,
the lower layers
hardly get trained


(sometimes they
get bigger & bigger,
then training diverges)

DIVING INTO THE MACHINE ROOM

- How training a neural network works
- **Activation functions**
- Faster optimizers
- Learning rate scheduling
- Regularization
- General suggestions


ACTIVATION FUNCTIONS

sigmoid




inspired
by biology

tanh



vanishing / dead gradients

ReLU




doesn't
saturate
here

works best
for shallow
networks


BETTER ACTIVATION FUNCTIONS

Leaky ReLU



$\max(\alpha z, z)$
 α hyperparameter
 $0 < \alpha < 1$


ELU



$$\begin{cases} \alpha(e^z - 1) & z < 0 \\ z & z \geq 0 \end{cases}$$

outperforms all
variants of ReLU

Swish



$z \cdot \text{sigmoid}(z)$
outperforms everything
for deep NNs

RECOMMENDATIONS

Try ReLU for shallow networks and Swish for deep networks

```
layer = Dense(100, activation="relu", kernel_initializer="he_normal")  
layer = Dense(100, activation="swish", kernel_initializer="he_normal")
```

how the values of weights
and biases are initialized

for reference


| Initialization method | Activation function |
|-----------------------|--|
| Glorot (default) | tanh, sigmoid, softmax |
| He | ReLU, Leaky ReLU, ELU, GELU, Swish, Mish |
| LeCun | SELU |

needs
to match activation
functions

DIVING INTO THE MACHINE ROOM

- How training a neural network works
- Activation functions
- **Faster optimizers**
- Learning rate scheduling
- Regularization
- General suggestions


FASTER OPTIMIZERS



#GradientDescent
UpdatingParameters



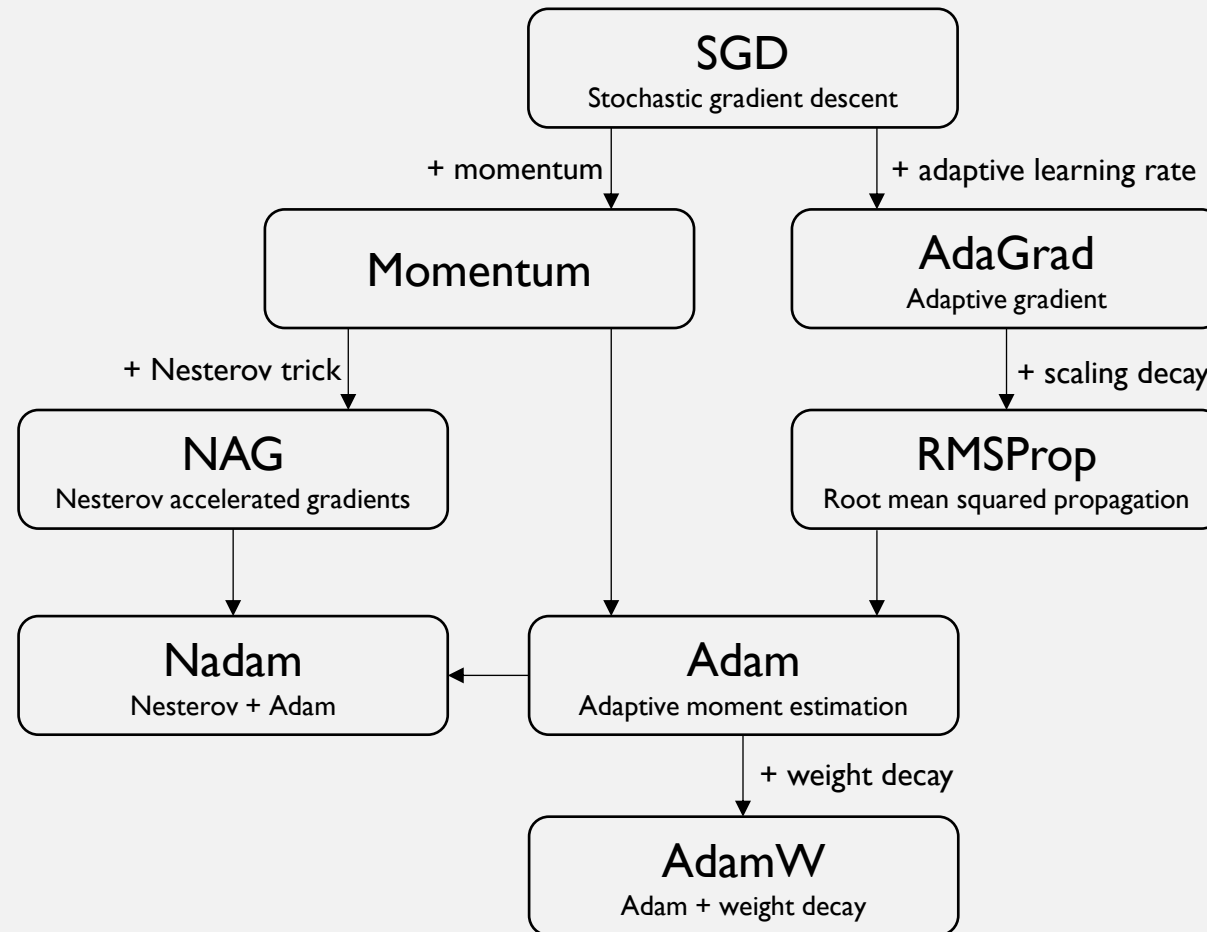
#GradientDescent
FinishingAnEpoch



#GradientDescent
WhenItConverges

The point is: Gradient descent can be painfully slow

FASTER OPTIMIZERS



momentum

Nesterov trick

adaptive learning rate

scaling decay

weight decay

MOMENTUM

"the step we just took probably wasn't a terrible idea"

Gradient descent


$$\theta \leftarrow \theta - \eta \nabla L(\theta)$$

Momentum

momentum

$$m \leftarrow \beta m - \eta \nabla L(\theta)$$

$$\theta \leftarrow \theta + m$$



```
optimizer = SGD(learning_rate=0.001, momentum=0.9)
```

The trick works because m generally points in the right direction, so the gradient here is slightly more accurate

THE NESTEROV TRICK

Momentum

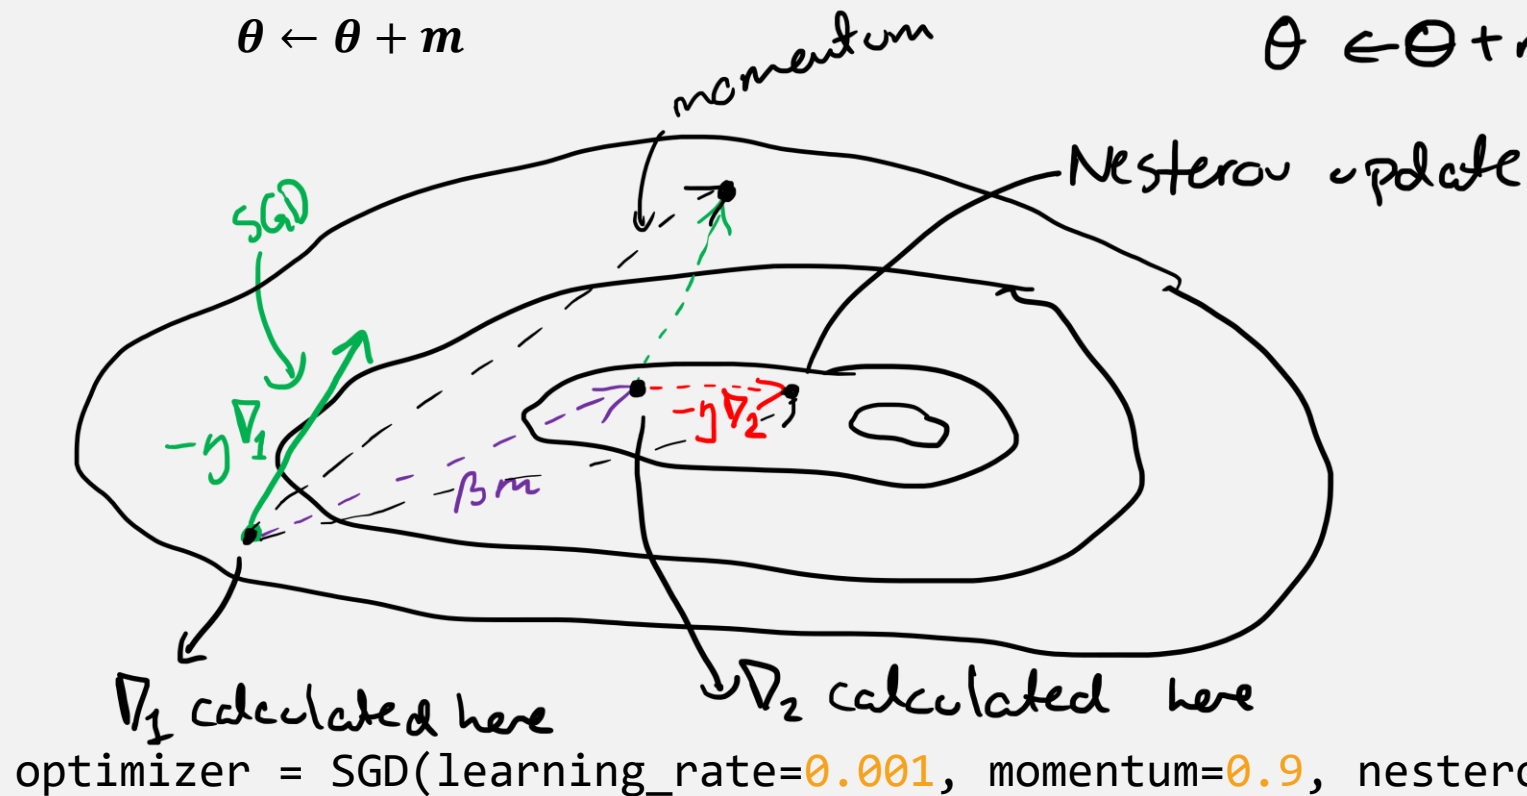
$$m \leftarrow \beta m - \eta \nabla L(\theta)$$

$$\theta \leftarrow \theta + m$$

Nesterov

$$m \leftarrow \beta m - \eta \nabla L(\theta + \beta m)$$


$$\theta \leftarrow \theta + m$$



momentum
Nesterov trick
adaptive learning rate
scaling decay
weight decay

ADAPTIVE LEARNING RATES

Observation: Loss functions often resemble elongated bowls



momentum
Nesterov trick
adaptive learning rate
scaling decay
weight decay

SCALING DECAY

add decay to s so it doesn't explode

AdaGrad

$$s \leftarrow s + \nabla L(\theta) \otimes \nabla L(\theta)$$

$$\theta \leftarrow \theta - \eta \nabla L(\theta) \oslash \sqrt{s + \varepsilon}$$

$\nearrow \sim 0.9$
RMSProp

$$s \leftarrow \rho s + (1 - \rho) \nabla L(\theta) \otimes \nabla L(\theta)$$

\rightarrow same thing

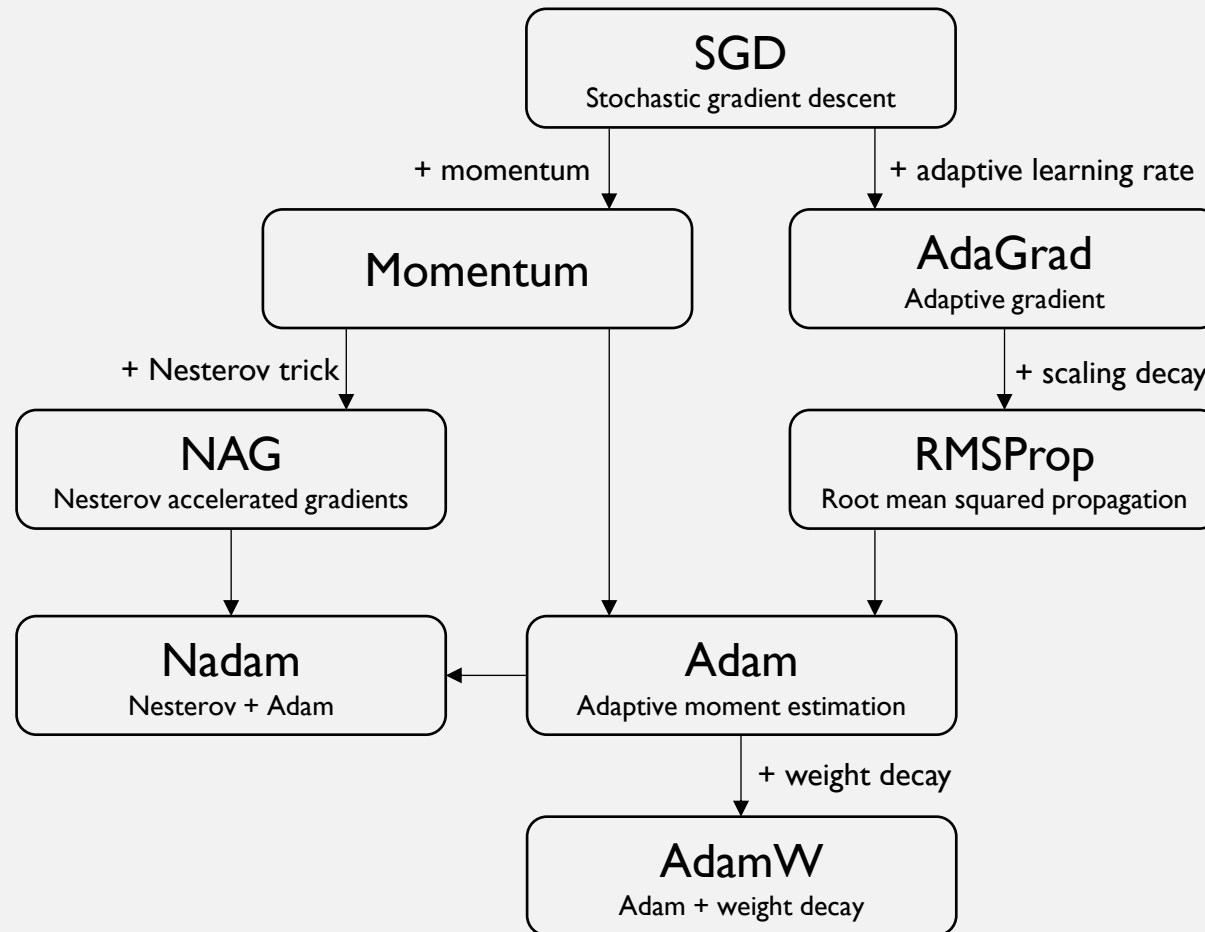
momentum
Nesterov trick
adaptive learning rate
scaling decay
weight decay

WEIGHT DECAY

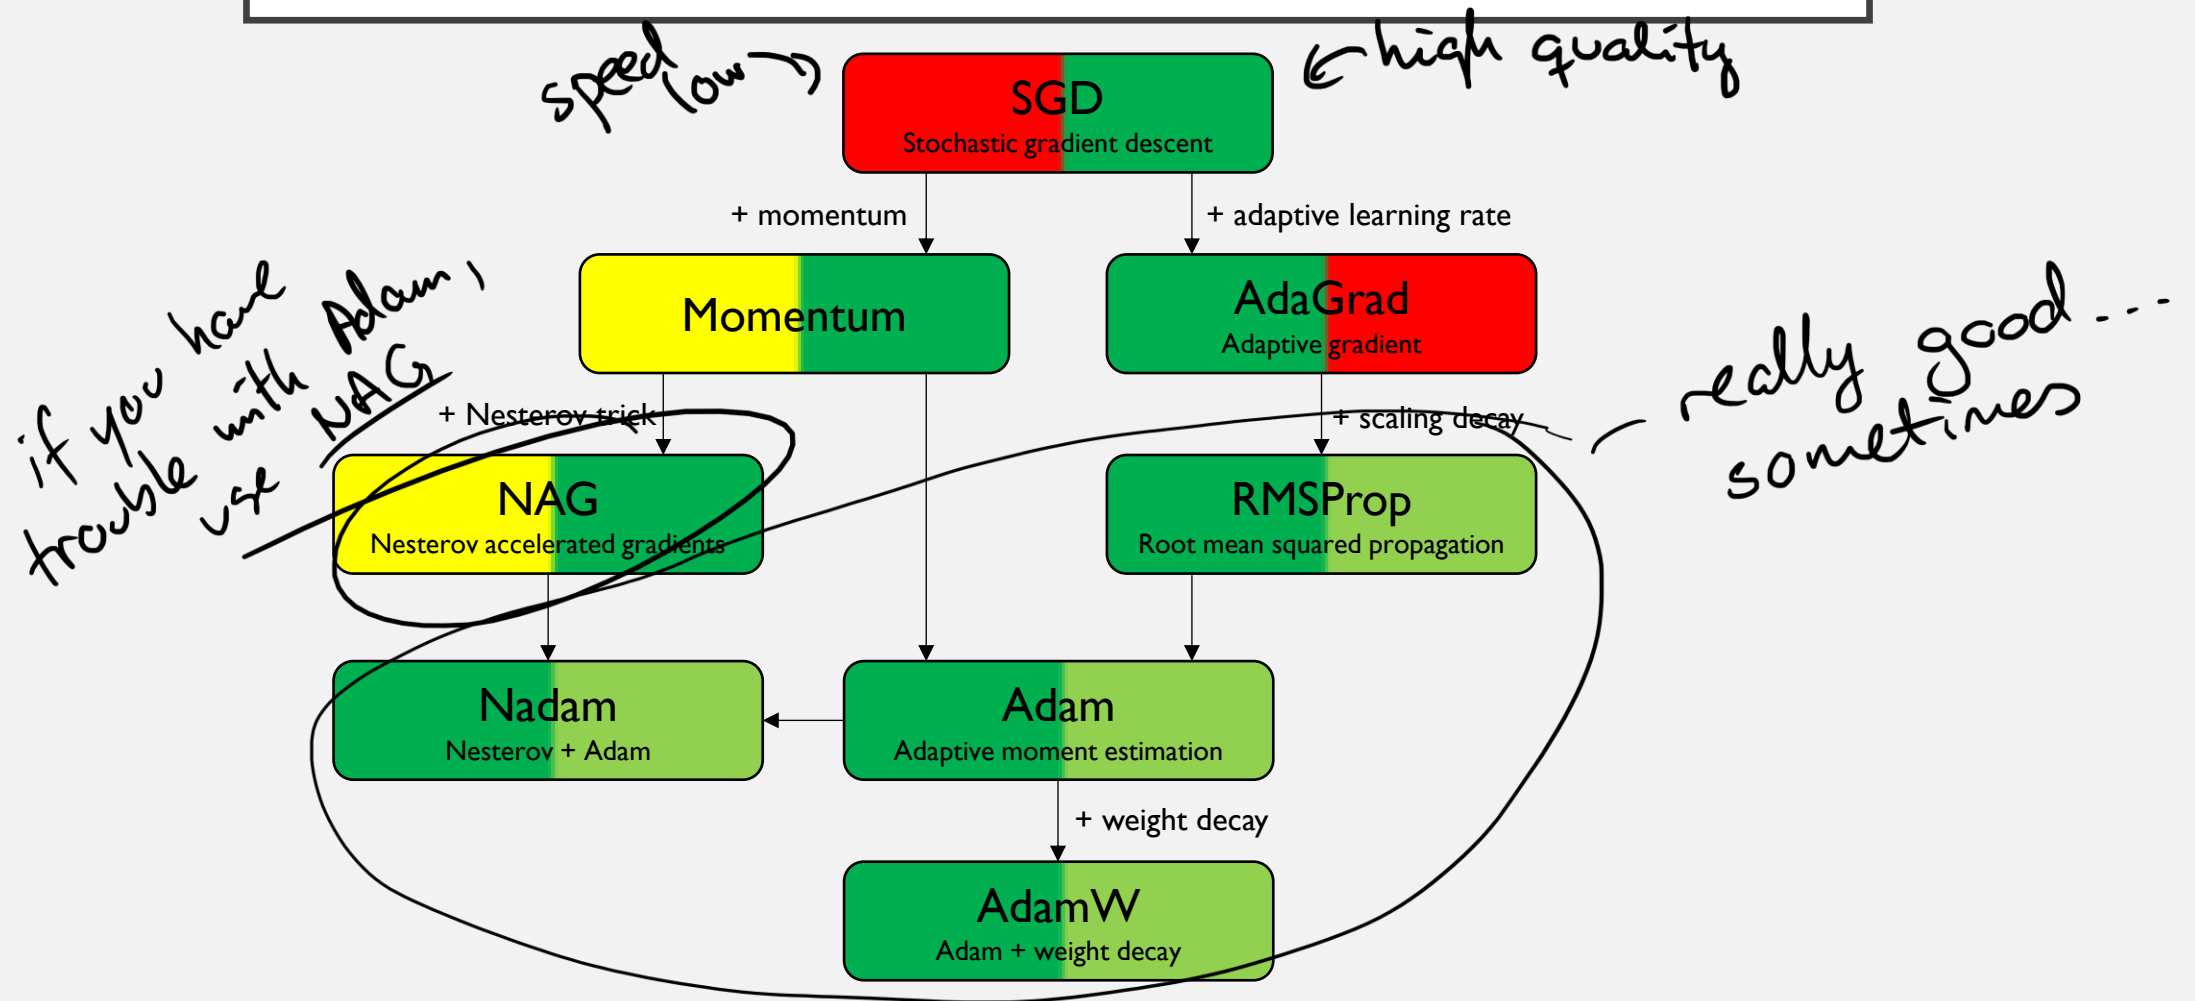
At each training step, multiply
all weights by, say 0.99

→ built-in regularization

RECOMMENDATIONS



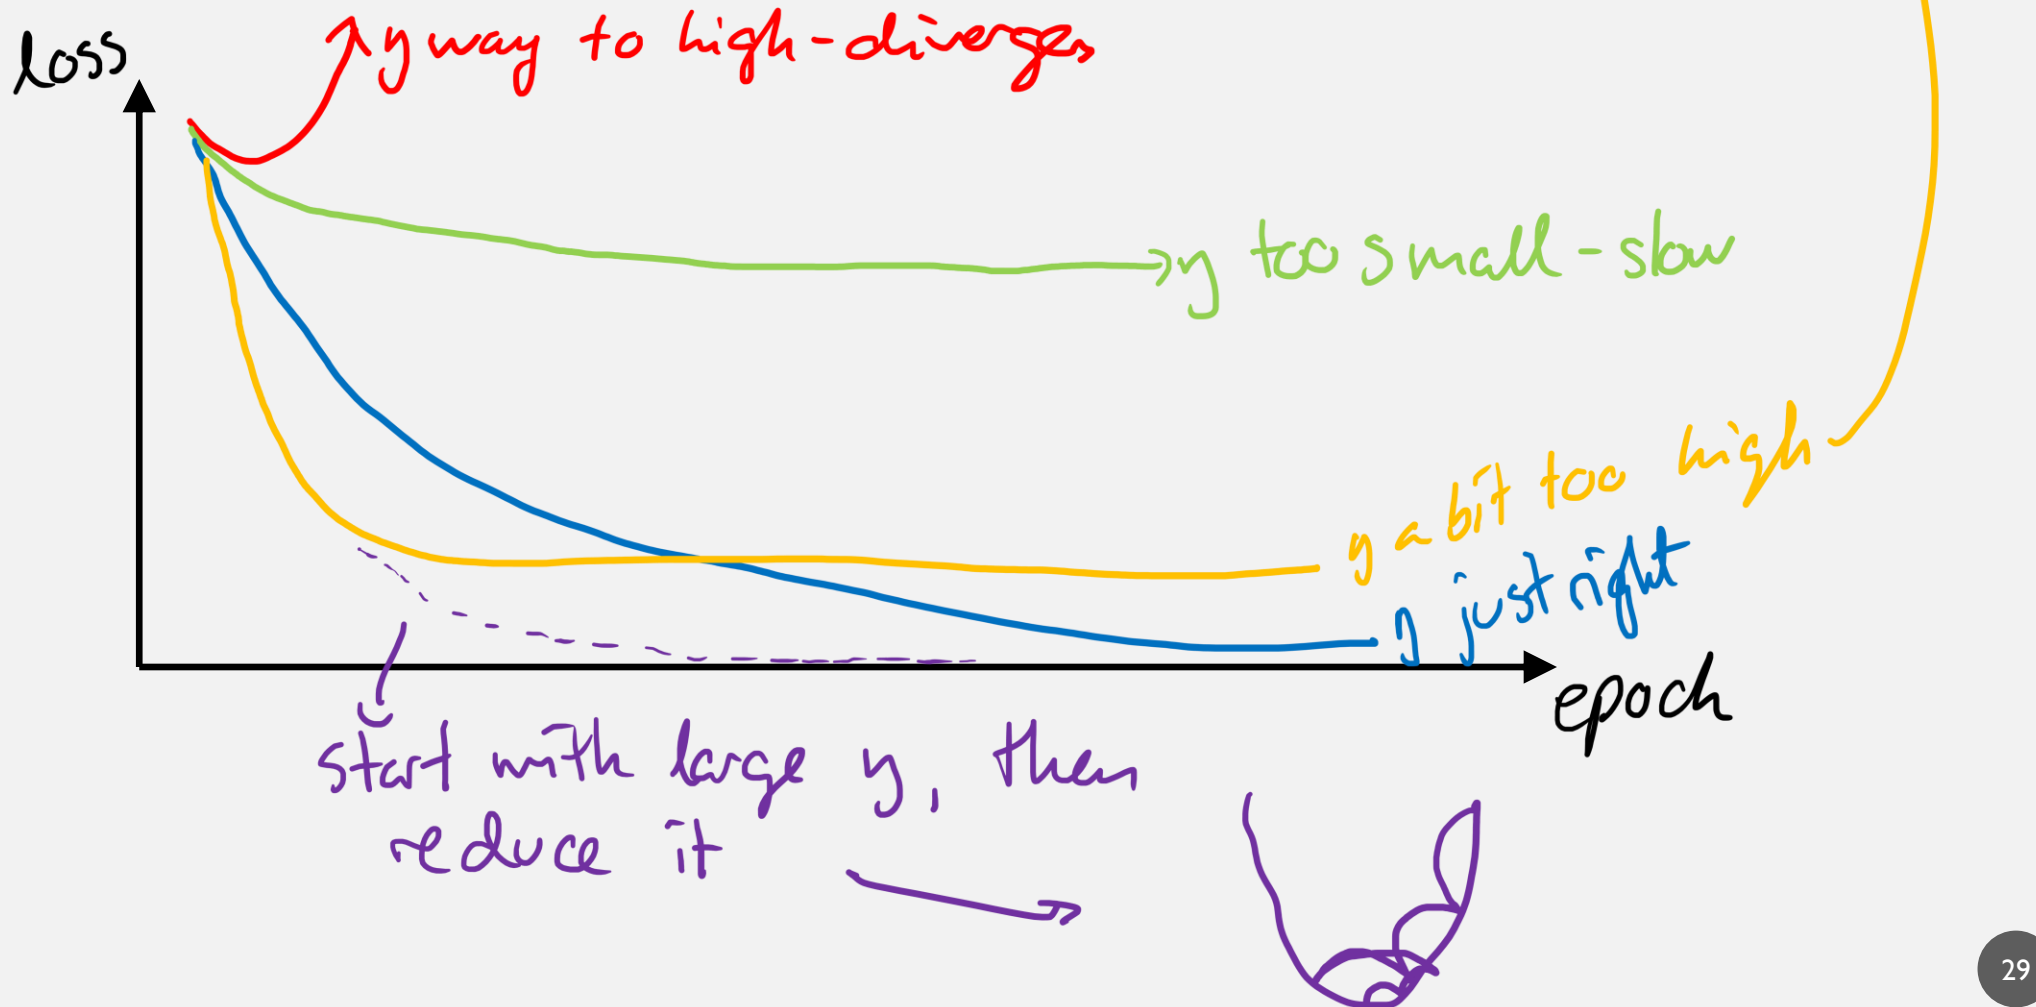
RECOMMENDATIONS



DIVING INTO THE MACHINE ROOM

- How training a neural network works
- Activation functions
- Faster optimizers
- **Learning rate scheduling**
- Regularization
- General suggestions

LEARNING RATE SCHEDULING



LEARNING RATE SCHEDULING

Power

$$\eta(t) = \frac{\eta_0}{1+t/s}$$

iteration #
after s steps: $\frac{\eta_0}{2}$
2s $\frac{\eta_0}{3}$
3s $\frac{\eta_0}{4}$

Exponential

$$\eta(t) = \eta_0 \cdot 0.1^{t/s}$$

reduces by a
factor of 10
every s steps

Piecewise constant

$$\eta = 0.1 \text{ for } 5 \text{ epochs}$$

$$\eta = 0.01 \text{ for } 10$$

$$0.001 \text{ for } 50$$

...

Performance

measure valid.
error every
N steps, divide
 η by λ when
error stops
dropping

All of them work pretty well – performance scheduling is a good default

CODING A LEARNING RATE SCHEDULE

```
#power scheduling
optimizer = SGD(learning_rate=0.01, decay=1e-4)

#exponential scheduling
def exp_decay_fn(epoch):
    return 0.01 * 0.1 ** (epoch/20)

lr_scheduler = tf.keras.callbacks.LearningRateScheduler(exp_decay_fn)
history = model.fit(..., callbacks = [lr_scheduler])

#piecewise constant scheduling
def piece_fn(epoch):
    if epoch < 5:
        return 0.01
    elif epoch < 15:
        [...]
    and so on - use LearningRateScheduler callback with piece_fn

#performance scheduling
lr_scheduler = tf.keras.callbacks.ReduceLROnPlateau(factor = 0.5, patience = 5)
history = model.fit(..., callbacks = [lr_scheduler])
```

DIVING INTO THE MACHINE ROOM

- How training a neural network works
- Activation functions
- Faster optimizers
- Learning rate scheduling
- **Regularization**
- General suggestions

L1 AND L2 REGULARIZATION

tools to avoid overfitting

Remember Lasso and Ridge regression?

penalty for large coefficients

$$\hookrightarrow \alpha \sum_i \theta_i^2 \quad (L2)$$

$$\alpha \sum_i |\theta_i| \quad (L1)$$


never use
with types
Adam - use AdamW
instead

```
layer = Dense(..., kernel_regularizer=tf.keras.regularizers.l2(0.01))
```

DROPOUT REGULARIZATION

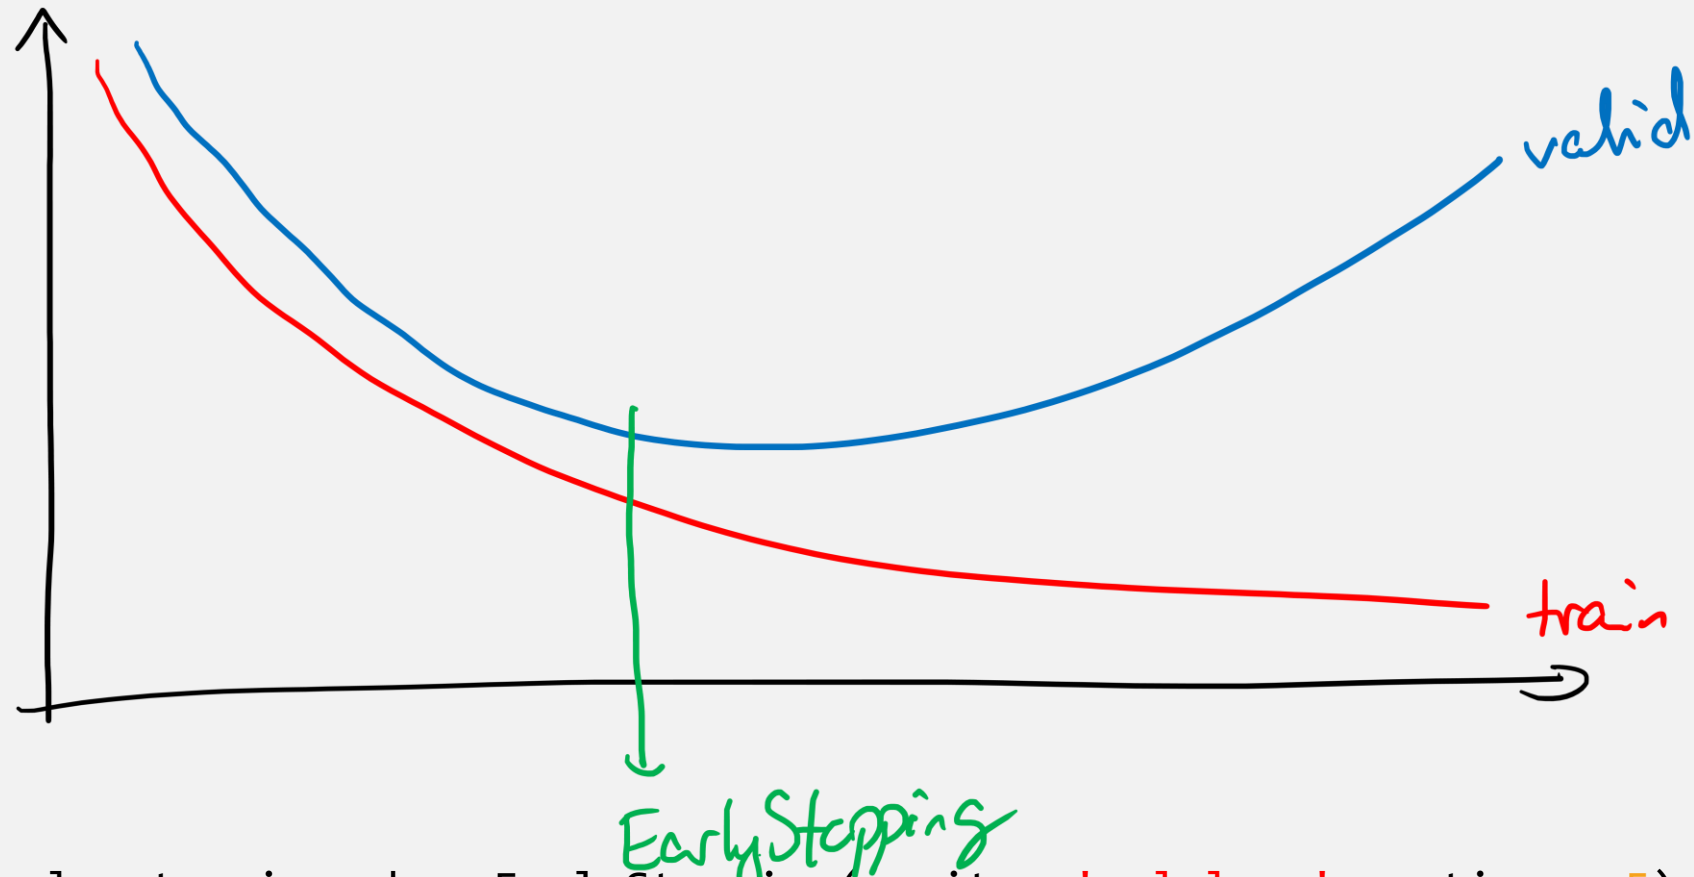
At every training step, every neuron has a probability p of being entirely ignored!

So the model can't
rely too much on
any particular neuron,
making it more robust



```
tf.keras.layers.Dropout(0.2)
```

BUT THE BEST WAY TO PREVENT
OVERFITTING IS USUALLY ...



```
early_stopping_cb = EarlyStopping(monitor='val_loss', patience=5)  
history = model.fit(..., callbacks = [early_stopping_cb])
```

DIVING INTO THE MACHINE ROOM


- How training a neural network works
- Activation functions
- Faster optimizers
- Learning rate scheduling
- Regularization
- **General suggestions**

GENERAL SUGGESTIONS

| Hyperparameter | Recommendations |
|------------------------|--|
| Activation function | ReLU if shallow. Swish if deep (Leaky ReLU as a faster alternative). |
| Optimizer | Start with Adam or AdamW, but switch to NAG if it doesn't work out. Never use SGD or AdaGrad. |
| Learning rate schedule | Performance scheduling is pretty good and requires very little hyperparameter tuning, but others may do better if you do it right. |
| Regularization | EarlyStopping and weight decay will usually do the trick, but look at others if you can't get rid of overfitting. Never use L1/L2 with Adam-type optimizers. |

TWO TASKS

**Find your neural network
from last week**



and make it better now that
you are cleverer

Scan this QR code



and tell me about something
you are still unsure about

You have 20 minutes