

GENERATIVE ADVERSARIAL NETWORKS


Lecture 5

MAL2, Spring 2025


GENERATIVE ADVERSARIAL NETWORKS

- Two neural networks fighting
- The troubles of training
- Deep convolutional GANs
- Conditional GANs
- Text-to-image generation

TWO NEURAL NETWORKS FIGHTING

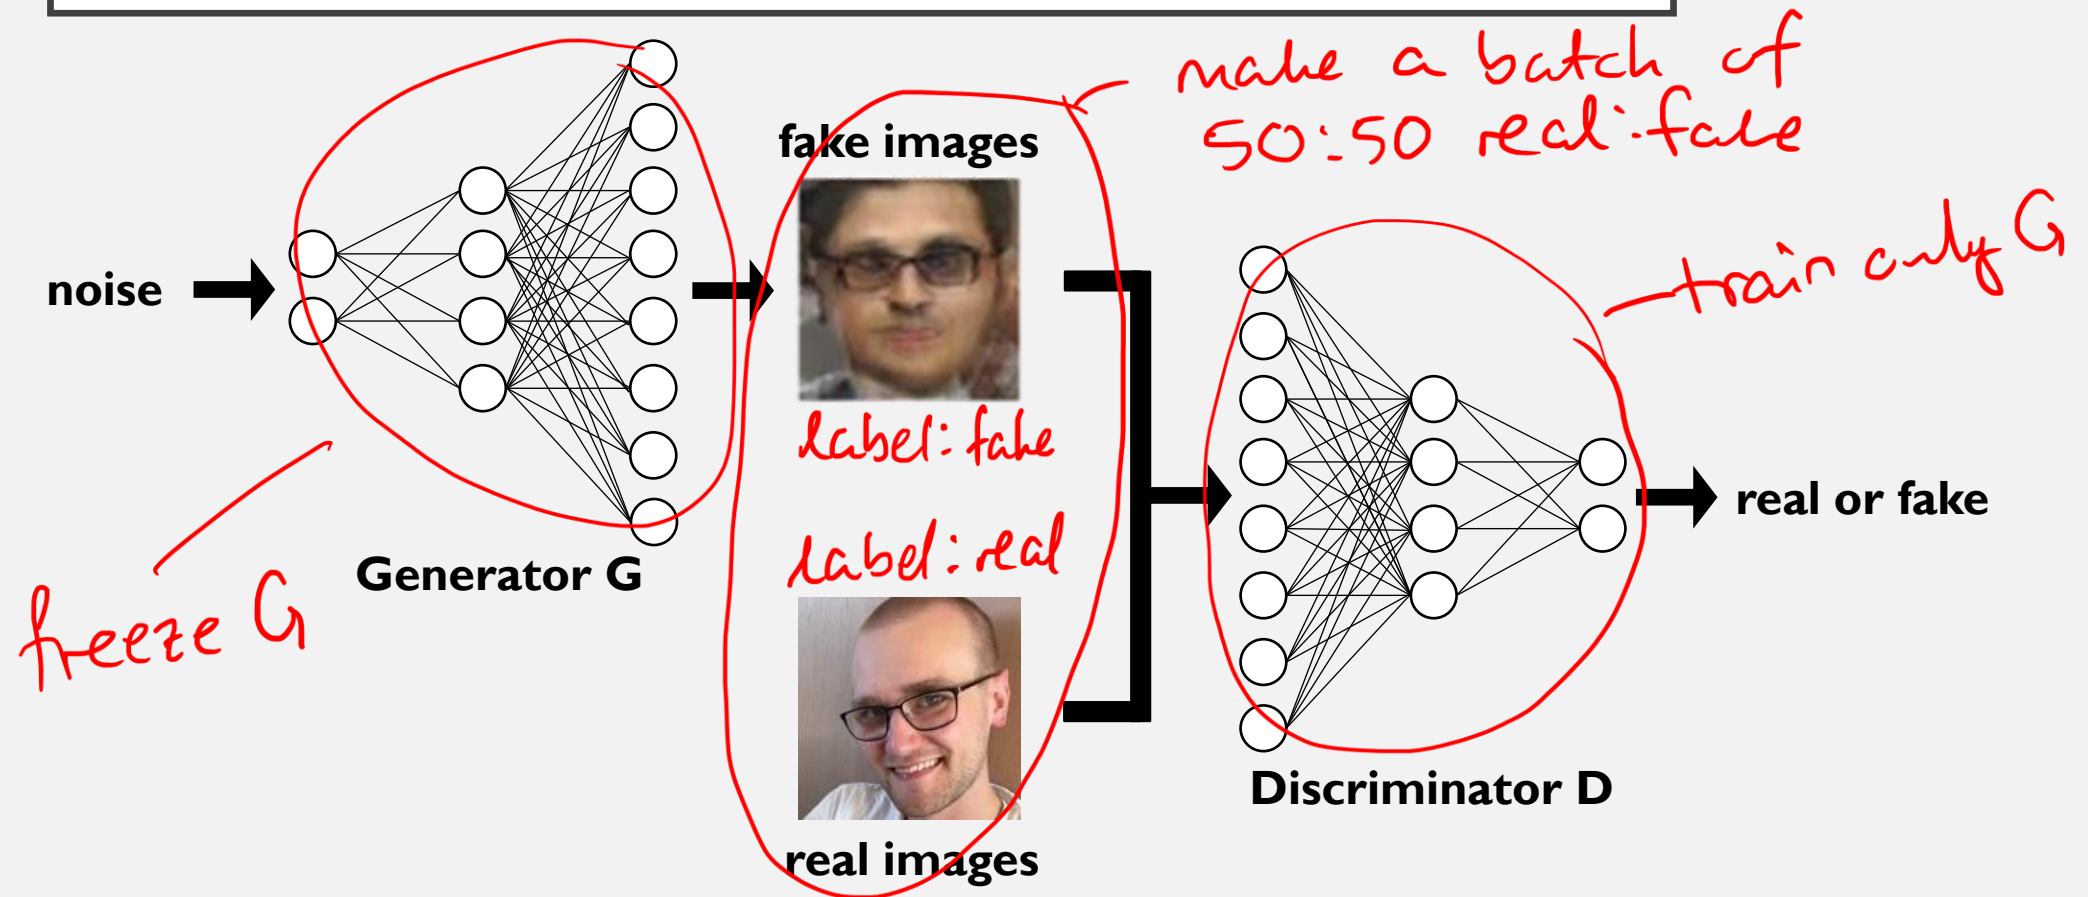


TWO NEURAL NETWORKS FIGHTING



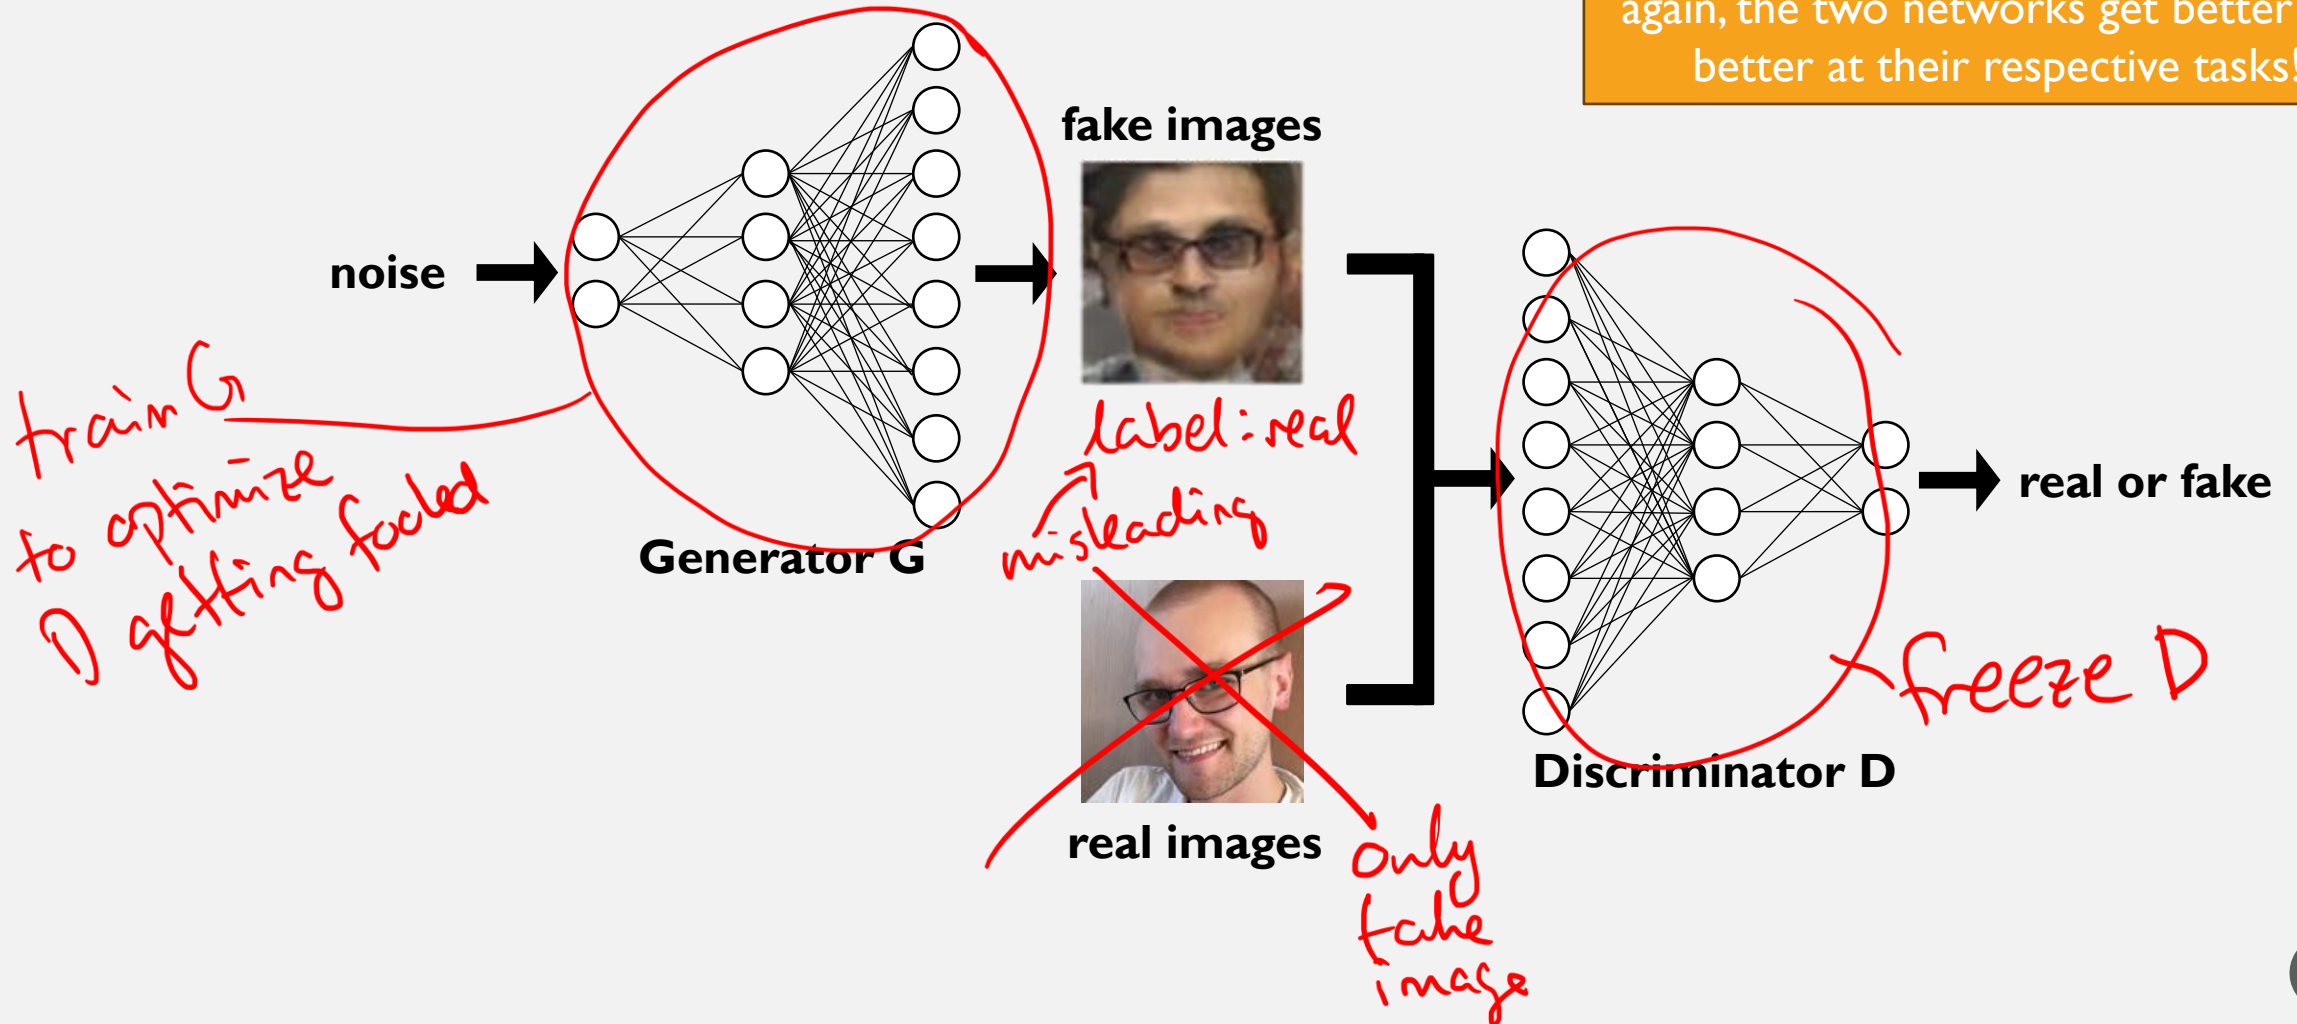
Training goals: D: get good at
classifying real/fake image
G: fool the discriminator

PHASE I: TRAIN THE DISCRIMINATOR




PHASE II: TRAIN THE GENERATOR

By repeating these two phases again and again, the two networks get better and better at their respective tasks!



LET'S DO IT



LET'S DO IT

Take the **GAN** we just made



and experiment with

- *the network architectures*
- *the learning rates*
- *the optimizers*
- *or something else*


Can you make training diverge?

You have 15 minutes

GENERATIVE ADVERSARIAL NETWORKS


- Two neural networks fighting
- The troubles of training
- Deep convolutional GANs
- Conditional GANs
- Text-to-image generation

MODE COLLAPSE




when G's output
becomes very
homogenous


MODE COLLAPSE




MODE COLLAPSE




MODE COLLAPSE




MODE COLLAPSE




MODE COLLAPSE




MODE COLLAPSE



MODE COLLAPSE



MODE COLLAPSE



SOLUTIONS TO MODE COLLAPSE

- Experience replay
 - Store images produced by the generator in a replay buffer
- Mini-batch discrimination
 - Measure similarity across a generated batch and give this information to D so that it can reject any homogeneous batch

OTHER CONVERGENCE CHALLENGES I

Too easy for D so G never learns

Epoch 1/20

1719/1719 ————— 21s 7ms/step - d_loss: 0.6214 - g_loss: 3.7936

Epoch 2/20

1719/1719 ————— 3s 2ms/step - d_loss: 0.4410 - g_loss: 2.6734

Epoch 3/20

1719/1719 ————— 3s 2ms/step - d_loss: 0.4148 - g_loss: 6.4556

Epoch 4/20

1719/1719 ————— 3s 2ms/step - d_loss: 2.1814e-05 - g_loss: 10.0660

Confuse the discriminator a bit:

```
labels += 0.05 * tf.random.uniform(tf.shape(labels))
```

OTHER CONVERGENCE CHALLENGES II

Epoch 1/200

36/36 ————— 6s 66ms/step - accuracy: 0.7405 - d_loss: 3.6339 - g_loss: 14.3425

poor D on real/fake

D ok, G bad

Epoch 9/200

36/36 ————— 0s 2ms/step - accuracy: 0.5009 - d_loss: 0.7056 - g_loss: 0.8502

*convergence
fleeing* →

perfect G, D forced to guess

Epoch 42/200

36/36 ————— 0s 2ms/step - accuracy: 0.6713 - d_loss: 0.6404 - g_loss: 1.9847

trying to make D better, so G becomes worse

Epoch 50/200

36/36 ————— 0s 2ms/step - accuracy: 0.9703 - d_loss: 0.0529 - g_loss: 7.7887

perfect D, G doesn't know what to do

Epoch 107/200

36/36 ————— 0s 2ms/step - accuracy: 1.0000 - d_loss: 7.0126e-07 - g_loss: 46.5934

THE TROUBLES OF TRAINING



In general, GANs are unstable and you may have to spend more effort than ever fine-tuning hyperparameters!


GENERATIVE ADVERSARIAL NETWORKS

- Two neural networks fighting
- The troubles of training
- **Deep convolutional GANs**
- Conditional GANs
- Text-to-image generation

DEEP CONVOLUTIONAL GANS

If we are working with images, surely we should use CNNs?

↓
Yes, but training is even more unstable



Guidelines

- Replace any pooling layers with strided convolutions in D and transposed convolutions in G.
- Use batch normalization in both G and D, except in the output layer of G and the input layer of D.
- Remove fully connected hidden layers for deeper architectures.
- Use ReLU in G for all layers except the output layer, which should use tanh.
- Use leaky ReLU in all layers of D.
- Don't trust the guidelines too much.



Even though the latent vectors are random, they are trained to represent meaningful features!

LATENT VECTOR ARITHMETIC



GENERATIVE ADVERSARIAL NETWORKS

- Two neural networks fighting
- The troubles of training
- Deep convolutional GANs
- **Conditional GANs**
- Text-to-image generation

CONDITIONAL GANs (cGANs)

$[0.21, 0.31, -0.49, 0.30, -0.11]$

noise
label

Generator

fake images
label

Discriminator

real or fake

$[0, 0, 0, 0, 0, 1, 0, 0, 0, 0]$

0 1 2 3 4 5

ONE-HOT
ENCODING

$\boxed{5} \Rightarrow \text{real}$
"5"

$\boxed{\text{scribble}} \Rightarrow \text{fake}$
"5"

$\boxed{6} \Rightarrow \text{fake}$
"5"

CONDITIONAL GANs (cGANs)

With two distinct inputs, the model is not just a sequence of layers, so Sequential is useless.

```
def build_discriminator():  
    → image_input = Input(shape=(img_dim, img_dim))  
    → label_input = Input(shape=(num_classes,))  
    flat_image = Flatten()(image_input)  
    concat = Concatenate()([flat_image, label_input])  
  
    { x = SomeLayer(...)(concat)  
      y = SomeOtherLayer(...)(x)  
    } real_or_fake = Dense(1, activation="sigmoid")(y)  
  
    return tf.keras.Model([image_input, label_input], real_or_fake)  
  
discriminator = build_discriminator()
```

and similar for the generator

PORTFOLIO ASSIGNMENT 2




GENERATIVE ADVERSARIAL NETWORKS

- Two neural networks fighting
- The troubles of training
- Deep convolutional GANs
- Conditional GANs
- Text-to-image generation

TEXT-TO-IMAGE GENERATION

Great idea for a challenging final project!



YOUR TICKET OUT THE DOOR

Scan this QR code



and tell me about something
you are still unsure about